


HEARTS

Human Expertise for AI Red-Teaming and Scalable Evaluation

DATE	TIME	FORMAT	ATTENDEES
April 14, 2026	2:15–6:00 CET	In person · 3 hrs	40–60



human · AI · evaluation

WHAT IS AI RED TEAMING?

Stress-testing AI before it reaches the world

Red teaming means adversarially probing an AI system to find risks before they cause real-life harms. Red-teamers intentionally surface unsafe outputs, biases, and other types of harms, drawing from military and cybersecurity traditions where simulated attacks reveal gaps early.

It's now a practice increasingly used at major tech companies and has even been encouraged by governments and public challenges like DEFCON 2023 and NIST's ARIA program. The challenge: there's an increasing need for red-teaming, but it's unclear how we can approach this challenge of scale while still retaining quality and centering the underlying human expertise necessary.

- Adversarial testing
- AI safety
- Risk identification
- Human evaluation

PROGRAMME

- 2:15 CET**
Welcome and introductions
Overview of AI red teaming and a hands-on red teaming exercise with insight from red-teamer Bolor-Erdene Jagdagdorj
- 2:45 CET**
Session 1 — Envisioning at scale
Breakout groups on automation, coverage, and expert oversight
- 3:45 CET**
Break and networking
- 4:15 CET**
Keynote + Q&A
Alex Taylor and Ding Wang on AI as an ecosystem
- 4:45 CET**
Session 2 — Zoomed-out discussion
Broader conversation on the future of human-centered AI evaluation
- 5:45 CET**
Closing reflection
Takeaways, next steps, setting up networks

THREE THEMES

Automation levels
When agents help and when humans must lead the evaluation

Evaluator well-being
Preventing harms from repeated exposure to harmful content

Tooling gaps
Scenario libraries, annotation tools, and triage pipelines

WHAT YOU GET

A shared output on key challenges and guiding principles for scaling AI red teaming

Connections with practitioners across academia, industry, and nonprofits

Pathways toward future publications, white papers, and grant proposals

ORGANISED BY

CMU · U. of Edinburgh ·
Google · Microsoft · Humane
Intelligence

PARTICIPANTS

40–60
academics, red teamers,
policymakers

SECTORS

3
academia · industry · nonprofits

SUPPORTED BY

NSF GRFP (DGE2140739) and
NIST (60NANB24D231)