

Statement of Interest

Who Gets to Evaluate? Distributed Accountability in Red Teaming

David Hartmann

TU Berlin

Berlin, Germany

Weizenbaum Institute for the Networked Society

Berlin, Germany

d.hartmann@tu-berlin.de

1 Introduction

I am a PhD student at the Technical University of Berlin and the Weizenbaum Institute for the Networked Society, working at the intersection of critical data studies, machine learning (ML), and human-computer interaction (HCI). My core interest is *distributed accountability and evaluation structures*: who can produce credible evidence about deployed AI systems, under what access conditions, with which measurement standards, and for whose benefit. This perspective treats evaluation not as a neutral technical afterthought but as an infrastructure—a set of methods, documentation practices, labour arrangements, and institutional incentives that shape what counts as harm, what becomes visible, and what remains unaddressed [13, 28, 37].

This matters especially for red-teaming and scalable evaluation of generative AI. Many central risk concepts — e.g., “toxicity,” “hate speech,” “deception,” “misinformation,” or “sycophancy” — are not merely engineering targets but socially contested and context-dependent categories [3, 9, 33]. In such settings, validity hinges on explicit choices about *what* we measure and *how* we map complex social phenomena onto operational indicators [2, 13, 16]. At scale, these choices are often delegated to benchmarks, rating rubrics, scenario libraries, or automated guardrail models—even though benchmarked performance can be unstable over time and sensitive to shifts in prompts, versions, and usage contexts [8, 23].

Reporting and documentation practices can make evaluation choices inspectable, but only if they are tied to concrete accountability pathways and independent scrutiny [7, 12, 27]. Otherwise, evaluation risks drifting into abstraction: technical proxies are treated as if they exhaust socio-technical harms, while the most consequential questions (who benefits, who bears the costs, who gets standing as an evaluator) remain implicit [9, 26].

My work, therefore, asks: **what would it mean to build evaluation infrastructures that are (i) methodologically rigorous, (ii) oriented toward affected users and marginalised groups, and (iii) robust against corporate capture?** This is the lens through which I am excited about HEARTS: centring human expertise and well-being is not an “ethics add-on,” but a prerequisite for valid and accountable red teaming at scale [15, 22, 25].

2 My related published and current work

Across my projects, I try to connect three layers: (1) technical methods for auditing black-box systems under access constraints, (2) HCI accounts of how evaluation expertise and workflow support are organised, and (3) governance questions about who can evaluate, under what rules, and with what evidentiary standing [7, 25, 28].

Auditing content moderation APIs (CHI 2025). In *Lost in Moderation* [18], we audited and red-teamed five widely used commercial content moderation APIs at scale. Analysing five million queries based on four datasets, we tested group-targeted hate speech and linguistic variation and found systematic patterns of *over-moderation* (silencing counterspeech, reclaimed slurs, or discussions about marginalised groups) and *under-moderation* (missing implicit, coded, or euphemistic hate, with especially concerning gaps for LGBTQIA+ targeted content). The contribution is not only empirical but methodological: a reproducible black-box auditing approach that combines careful dataset design with stress tests as well as qualitative evaluations of model failures for linguistic variation and conceptually difficult cases.

This work is aligned with longstanding concerns that hate speech detection and “toxicity labelling can encode racialised and dialect-related biases, and that evaluation must be attentive to sociolinguistic context [31–33]. It also connects to the broader debate about fairness metrics and their failure modes: threshold-agnostic bias metrics and AUC-based views can be useful, but have important limitations and can mislead if treated as stand-ins for downstream harms [4, 5].

Query-efficient active fairness auditing of black-box LLMs (preprint online, submitted). A recurring barrier to independent evaluation is the cost structure of APIs: rate limits, paywalls, and limited observability make rigorous audits expensive [6, 40]. My current technical work addresses this by developing a query-efficient *active* auditing method for black-box language models [19]. The core idea is to maintain uncertainty about a fairness-relevant metric (e.g., group differences in ROC–AUC) via a family of surrogate hypotheses consistent with observed black-box outputs, and to choose new queries expected to shrink uncertainty most. This builds on active auditing and active learning foundations [1, 34, 38] and connects to recent work emphasising that credible audits require explicit evidence standards, access descriptions, and uncertainty reporting [7, 18]. Practically, this line of work is motivated by the reality that most third-party auditors face bounded budgets and must treat evaluation itself as a resource allocation problem [36, 40] and will be, hopefully, followed-up by HCI-focused research on an interface showcasing such uncertainty of different metrics that can be adjusted thorough querying the black-box model by auditors.

Civil society evaluation and a participatory audit toolkit (ongoing). In parallel, I work closely with my research group lead, Milagros Miceli and HCI scholar Goda Klumbyte, on civil society evaluation practices in the EU and a participatory audit toolkit. The starting

point is that transparency alone is insufficient for accountability: people need the ability to produce, interpret, contest, and act on evidence, especially under constrained access conditions [7, 28]. This connects to HCI research on user- and community-driven auditing, including how everyday and scaffolded auditing can surface harms missed by formal evaluations [10, 11, 35].

Depending on the participatory process, in which civil society itself decides the process and outcome, we hope to translate these insights into practical resources for civil society auditing, such as a platform or wiki that describes clusters of strategies and methods for understanding, auditing, and contesting AI systems from a CSO perspective. The clusters identified so far are: (1) legal contestation via FOIA and GDPR complaints, (2) socio-technical evaluations that combine interviews with technical evaluations, (3) infrastructural evaluations which focus on how a product is embedded in infrastructures and what kind of risks arise from that and (4) the technical analysis, doing classic fairness and platform auditing. Our work so far also aligns with practitioner-centred work showing that real-world fairness and auditing practice needs socio-technical support beyond metrics—including processes, tooling, and institutional pathways [20, 22, 24, 25].

Honeypot crawler audit of “consent” signals (ongoing). Another ongoing project tests whether generative AI crawlers and agents respect web-based consent signals (e.g., robots.txt files) using honeypot researcher websites. While not “red teaming” in the narrow security sense, it is an adversarial infrastructure evaluation as we probe compliance under realistic deployment conditions and highlight how contested norms (what counts as consent, what counts as public data) intersect with corporate incentives.

Regulatory and ecosystem work. Across these threads, I argue for an EU AI audit ecosystem that treats independent evaluation as a first-class accountability mechanism, including civil society and academic auditors as participants with standing [17]. This complements HEARTS’ aims by emphasizing that scalable red teaming must be supported by durable institutions and access regimes, not only better tools [14, 40].

3 My position on red teaming: evaluation crisis, corporate capture, and labor

I strongly support rigorous methods to identify and mitigate risks in deployed AI systems. At the same time, I find the current framing of “red teaming” sometimes problematic—not because adversarial testing is unhelpful, but because the term often imports a particular *security mindset* and *organizational logic* that can narrow what gets evaluated and why [15].

1) Scientific invalidity risks: from exploration to evidence. Red teaming is often exploratory: generate adversarial prompts, find failures, patch. That exploration is valuable, but it does not automatically yield valid evidence about prevalence, impacts on affected groups, or real-world consequences. Without explicit constructs and validity arguments, red teaming can blur the line between “we found a scary example” and “we understand the risk profile of our system” [2, 13, 16]. This is especially acute for contested concepts like hate speech, where disagreement is not noise but part of the phenomenon [3, 9, 33]. Scaling evaluation, therefore, requires more

than bigger prompt sets. For me, it requires disciplined construct definitions, uncertainty reporting, and claims that are proportional to the evidence [12, 29].

2) Corporate capture risks: red teaming as ritual or “security theater.” When red teaming is conducted primarily by providers or tightly controlled partners, it can become a legitimacy practice: producing reassuring narratives without enabling independent scrutiny. Recent critiques warn that red teaming can drift toward “security theater” if it emphasizes optics, secrecy, or patch counts over transparent assessment [15]. This risk is amplified by access asymmetries: auditors are often restricted to black-box interfaces, while stronger forms of scrutiny depend on white-box and outside-the-box access (documentation, training and deployment artifacts) [6, 7]. Mapping work on who evaluates AI’s social impacts suggests that third-party evaluations remain systematically underprovided compared to developer-led assessments [30]—which is precisely why I treat independent evaluation as infrastructure to counter capture [28].

3) Labour and expertise: making evaluation work visible and valued. Scaling red teaming creates labour demands: scenario writing, annotation, adjudication, triage, and risk management. Yet this labour is often framed as interchangeable labeling, even though it requires contextual and domain expertise. In practice, evaluation support needs include workflow design, documentation, and organizational processes [20, 22, 24, 25]. In my group, working closely with Milagros Miceli, I engage with the data workers’ inquiry perspective to foreground the situated knowledge and constraints of people who sustain evaluation and data production pipelines. For HEARTS, I see evaluator well-being and compensation as *epistemic* issues and this that under-supported evaluators produce less reliable evidence.

4) Participation and accountability: including civil society and affected groups without offloading harm. Finally, I believe scalable evaluation must involve affected stakeholders (users, civil society organisations, community advocates)—not to outsource harm discovery, but to ensure that problem formulations, scenario libraries, and success criteria reflect those who bear consequences [10, 11, 35, 39]. Still, this remains an open design question: participation can impose burdens and create extractive dynamics if not designed with safeguards, power-sharing, and compensation [37, 39]. My participatory toolkit work aims to develop formats where civil society can shape evaluation priorities and interpretation, while minimizing direct exposure of affected individuals to harmful content unless they explicitly choose such roles.

In sum: I critique red teaming when it becomes (i) a narrow security ritual detached from validity, (ii) a provider-controlled legitimacy mechanism, and/or (iii) a labour system that hides expertise and harm. But I also see red teaming—broadly construed—as indispensable. We need:

- (1) methods to evaluate risks and harms *for users and affected groups* (not primarily for organizational reassurance) [2, 13],
- (2) independent evaluation as a standard infrastructure to counter corporate capture [28, 30, 40],

- (3) visibility and valuation of evaluation labour (monetarily and epistemically) [20, 22, 24], and
- (4) participatory safeguards that include civil society without shifting burdens onto those already harmed [10, 39].

4 What I can contribute and what interests me most

At HEARTS, I would contribute both (a) technical methods for constrained, query-efficient evaluation and (b) a socio-technical framing of red teaming as infrastructure and labour.

Concrete contributions (“Build” thread):

- **Uncertainty-aware red teaming templates:** reporting formats that separate exploratory findings from quantitative claims; explicit validity and scope statements; and “what would change our mind” checklists [12, 29].
- **Scenario library design for contested concepts:** guidelines for encoding context, target group, counterspeech, reclaimed language, and ambiguity [9, 31, 33].
- **Resource-constrained evaluation strategies:** practical approaches for auditors facing paywalls/rate limits (active sampling, disagreement targeting, stratified initialization), linked to governance needs for outsider oversight [1, 38, 40].
- **Labour-centred workflow patterns:** rotation and recovery practices, risk disclosure, and compensation norms for evaluators [21].

Key interests/questions (“Vision” thread):

- How can we define “scaling” as *improving epistemic coverage and validity* rather than merely increasing the number of adversarial prompts [2, 16]?
- What governance commitments (disclosure norms, independent access, documentation standards) would make red teaming outputs usable for public accountability rather than internal evaluation [7, 15]?
- How can we build evaluation infrastructures that are resilient to capture, enabling civil society and academic auditors to produce credible evidence [25, 28, 30]?

I see HEARTS as a chance to connect communities that are too often not speaking to one another: red team practitioners, HCI scholars of evaluation workflows and labour, and ML researchers building measurement methods. My goal is to help articulate a shared agenda where scalable red teaming becomes a humane, independent, and methodologically rigorous component of AI accountability infrastructures.

References

- [1] Anastasios N. Angelopoulos, Jacob Eisenstein, Jonathan Berant, Alekh Agarwal, and Adam Fisch. 2025. Cost-Optimal Active AI Model Evaluation. (2025). arXiv:2506.07949 [cs.LG] doi:10.48550/arXiv.2506.07949 arXiv preprint submitted 9 June 2025.
- [2] Andrew M. Bean, Ryan Othniel Kearns, Angelika Romanou, Franziska Sofia Hafner, Harry Mayne, Jan Batzner, Negar Foroutan, Chris Schmitz, Karolina Korgul, Hunar Batra, Oishi Deb, Emma Beharry, Cornelius Emde, Thomas Foster, Anna Gausen, Maria Grandury, Simeng Han, Valentin Hofmann, Lujain Ibrahim, Hazel Kim, Hannah Rose Kirk, Fangru Lin, Gabrielle Kaili-May Liu, Lennart Luettgau, Jabez Magomere, Jonathan Rystrom, Anna Sotnikova, Yushi Yang, Yilun Zhao, Adel Bibi, Antoine Bosselut, Ronald Clark, Arman Cohan, Jakob Nicolaus Foerster, Yarin Gal, Scott A. Hale, Inioluwa Deborah Raji, Christopher Summerfield, Philip Torr, Cozmin Ududec, Luc Rocher, and Adam Mahdi. 2025. Measuring what Matters: Construct Validity in Large Language Model Benchmarks. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems Datasets and Benchmarks Track*. <https://openreview.net/forum?id=mdA5lVvNcU>
- [3] Su Lin Blodgett, Solon Barocas, Hal Daumé III, and Hanna Wallach. 2020. Language (Technology) Is Power: A Critical Survey of “Bias” in NLP. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics* (Online). Association for Computational Linguistics, 5454–5476. doi:10.18653/v1/2020.acl-main.485
- [4] Daniel Borkan, Lucas Dixon, John Li, Jeffrey Sorensen, Nithum Thain, and Lucy Vasserman. 2019. Limitations of Pinned AUC for Measuring Unintended Bias. arXiv:1903.02088 [stat.ML] <https://arxiv.org/abs/1903.02088>
- [5] Daniel Borkan, Lucas Dixon, Jeffrey Sorensen, Nithum Thain, and Lucy Vasserman. 2019. Nuanced Metrics for Measuring Unintended Bias with Real Data for Text Classification. In *Companion Proceedings of The 2019 World Wide Web Conference* (San Francisco, USA) (WWW ’19). Association for Computing Machinery, New York, NY, USA, 491–500. doi:10.1145/3308560.3317593
- [6] Stephen Casper, Carson Ezell, Charlotte Siegmund, Noam Kolt, Taylor Lynn Curtis, Benjamin Bucknall, Andreas Haupt, Kevin Wei, Jérémy Scheurer, Marius Hobbhahn, Lee Sharkey, Satyapriya Krishna, Marvin Von Hagen, Silas Alberti, Alan Chan, Qinyi Sun, Michael Gerovitch, David Bau, Max Tegmark, David Krueger, and Dylan Hadfield-Menell. 2024. *Black-Box Access Is Insufficient for Rigorous AI Audits*. arXiv:2401.14446 [cs] doi:10.48550/arXiv.2401.14446
- [7] Sarah H. Cen and Rohan Alur. 2024. From Transparency to Accountability and Back: A Discussion of Access and Evidence in AI Auditing. In *Proceedings of the 4th ACM Conference on Equity and Access in Algorithms, Mechanisms, and Optimization* (San Luis Potosi Mexico, 2024-10-29). ACM, 1–14. doi:10.1145/3689904.3694711
- [8] Lingjiao Chen, Matei Zaharia, and James Zou. 2023. How Is ChatGPT’s Behavior Changing Over Time? *arXiv preprint arXiv:2307.09009* (2023).
- [9] Pieter Delobelle, Giuseppe Attanasio, Debora Nozza, Su Lin Blodgett, and Zeerak Talat. 2024. Metrics for What, Metrics for Whom: Assessing Actionability of Bias Evaluation Metrics in NLP. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, Singapore, 21669–21691. <https://aclanthology.org/2024.emnlp-main.1176> Joint first authors: Delobelle and Attanasio.
- [10] Wesley Hanwen Deng, Wang Claire, Howard Ziyu Han, Jason I. Hong, Kenneth Holstein, and Motahhare Eslami. 2025. WeAudit: Scaffolding User Auditors and AI Practitioners in Auditing Generative AI. *Proc. ACM Hum.-Comput. Interact.* 9, 7, Article CSCW521 (Oct. 2025), 35 pages. doi:10.1145/3757702
- [11] Alicia DeVos, Aditi Dhabalia, Hong Shen, Kenneth Holstein, and Motahhare Eslami. 2022. Toward User-Driven Algorithm Auditing: Investigating users’ strategies for uncovering harmful algorithmic behavior. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems* (New Orleans, LA, USA) (CHI ’22). Association for Computing Machinery, New York, NY, USA, Article 626, 19 pages. doi:10.1145/3491102.3517441
- [12] Jesse Dodge, Gabriel Ilharco, Roy Schwartz, Ali Farhadi, Hannaneh Hajishirzi, and Noah A. Smith. 2019. Show Your Work: Improved Reporting of Experimental Results. In *Proceedings of EMNLP-IJCNLP 2019*.
- [13] Maria Eriksson, Erasmo Purificato, Arman Noroozian, João Vinagre, Guillaume Chaslot, Emilia Gómez, and David Fernandez-Llorca. 2025. Can We Trust AI Benchmarks? An Interdisciplinary Review of Current Issues in AI Evaluation. In *Proceedings of the Eighth AAAI/ACM Conference on AI, Ethics, and Society (AI/ES 2025)*. Association for the Advancement of Artificial Intelligence, 850–.
- [14] European Commission. 2025. General-Purpose AI Code of Practice. <https://digital-strategy.ec.europa.eu/policies/contents-code-gpai> Accessed July 11, 2025.
- [15] Michael Feffer, Anusha Sinha, Wesley H. Deng, Zachary C. Lipton, and Hoda Heidari. 2025. *Red-Teaming for Generative AI: Silver Bullet or Security Theater?* AAAI Press, 421–437.
- [16] Sebastian Gehrmann, Elizabeth Clark, and Thibault Sellam. 2022. Repairing the Cracked Foundation: A Survey of Obstacles in Evaluation Practices for Generated Text. arXiv:2202.06935 [cs.CL] <https://arxiv.org/abs/2202.06935>
- [17] David Hartmann, José Renato Laranjeira De Pereira, Chiara Streitbürger, and Bettina Berendt. 2025. Addressing the Regulatory Gap: Moving towards an EU AI Audit Ecosystem beyond the AI Act by Including Civil Society. *AI and Ethics* (2025). doi:10.1007/s43681-024-00595-3
- [18] David Hartmann, Amin Oueslati, Dimitri Stauf, Lena Pohlmann, Simon Munzert, and Hendrik Heuer. 2025. Lost in Moderation: How Commercial Content Moderation APIs Over- and Under-Moderate Group-Targeted Hate Speech and Linguistic Variations. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems* (CHI ’25). doi:10.1145/3706598.3713998 Author’s version; accepted for CHI 2025.
- [19] David Hartmann, Lena Pohlmann, Lelia Hanslik, Noah Gießing, Bettina Berendt, and Pieter Delobelle. 2026. Audit Me If You Can: Query-Efficient Active Fairness Auditing of Black-Box LLMs. arXiv:2601.03087 [cs.LG] <https://arxiv.org/abs/2601.03087>
- [20] Kenneth Holstein, Jennifer Wortman Vaughan, Hal Daumé III, Miro Dudík, and Hanna Wallach. 2019. Improving Fairness in Machine Learning Systems: What Do Industry Practitioners Need?. In *CHI Conference on Human Factors in Computing*

- Systems (CHI '19). Association for Computing Machinery, New York, NY, USA, 16 pages. doi:10.1145/3290605.3300830
- [21] Airi Lampinen, Rob Comber, Sravya Chandhiramowuli, Naja Holten Møller, and Alex S. Taylor. 2024. The Work of AI: Mapping Human Labour in the AI Pipeline. In *Companion Publication of the 2024 Conference on Computer-Supported Cooperative Work and Social Computing* (San Jose, Costa Rica) (CSCW Companion '24). Association for Computing Machinery, New York, NY, USA, 728–731. doi:10.1145/3678884.3681825
 - [22] Tina B. Lassiter and Kenneth R. Fleischmann. 2024. “Something Fast and Cheap” or “A Core Element of Building Trust”? – AI Auditing Professionals’ Perspectives on Trust in AI. *Proceedings of the ACM on Human-Computer Interaction* 8, CSCW2 (November 2024), 424:1–424:22. doi:10.1145/3686963
 - [23] Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, Benjamin Newman, Binhang Yuan, Bobby Yan, Ce Zhang, Christian Cosgrove, Christopher D. Manning, Christopher Ré, Diana Acosta-Navas, Drew A. Hudson, Eric Zelikman, Esin Durmus, Faisal Ladhak, Frieda Rong, Hongyu Ren, Huaxiu Yao, Jue Wang, Keshav Santhanam, Laurel Orr, Lucia Zheng, Mert Yuksekgonul, Mirac Suzgun, Nathan Kim, Neel Guha, Niladri Chatterji, Omar Khattab, Peter Henderson, Qian Huang, Ryan Chi, Sang Michael Xie, Shibani Santurkar, Surya Ganguli, Tatsunori Hashimoto, Thomas Icard, Tianyi Zhang, Vishrav Chaudhary, William Wang, Xuechen Li, Yifan Mai, Yuhui Zhang, Yuta Koreeda, et al. 2023. Holistic Evaluation of Language Models. *Transactions on Machine Learning Research* (Aug. 2023). <https://openreview.net/forum?id=iO4LZibEqW> Published 23 Aug 2023. Also available as arXiv:2211.09110.
 - [24] Michael Madaio, Lisa Egede, Hariharan Subramonyam, Jennifer Wortman Vaughan, and Hanna Wallach. 2022. Assessing the Fairness of AI Systems: AI Practitioners’ Processes, Challenges, and Needs for Support. *Proceedings of the ACM on Human-Computer Interaction* 6, CSCW1 (2022), 52:1–52:26. doi:10.1145/3512899
 - [25] Victor Ojewale, Ryan Steed, Briana Vecchione, Abeba Birhane, and Inioluwa Deborah Raji. 2025. Towards AI Accountability Infrastructure: Gaps and Opportunities in AI Audit Tooling. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. Association for Computing Machinery, New York, NY, USA, Article 815, 29 pages. doi:10.1145/3706598.3713301
 - [26] Inioluwa Deborah Raji, Emily M. Bender, Amandalynne Paullada, Emily Denton, and Alex Hanna. 2021. AI and the Everything in the Whole Wide World Benchmark. In *Proceedings of the 35th Conference on Neural Information Processing Systems (NeurIPS)*. arXiv:2111.15366 [cs.LG] <https://arxiv.org/abs/2111.15366> Track on Datasets and Benchmarks.
 - [27] Inioluwa Deborah Raji, Andrew Smart, Rebecca N. White, Margaret Mitchell, Timnit Gebru, Ben Hutchinson, Jamila Smith-Loud, Daniel Theron, and Parker Barnes. 2020. Closing the AI accountability gap: defining an end-to-end framework for internal algorithmic auditing. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency (Barcelona, Spain) (FAT* '20)*. Association for Computing Machinery, New York, NY, USA, 33–44. doi:10.1145/3351095.3372873
 - [28] Inioluwa Deborah Raji, Peggy Xu, Colleen Honigsberg, and Daniel E. Ho. 2022. Outsider Oversight: Designing a Third Party Audit Ecosystem for AI Governance. <http://arxiv.org/abs/2206.04737> arXiv:2206.04737 [cs].
 - [29] Paulius Rauba, Qiyao Wei, and Mihaela van der Schaar. 2025. Statistical Hypothesis Testing for Auditing Robustness in Language Models. *arXiv preprint arXiv:2506.07947* (2025).
 - [30] Anka Reuel, Avijit Ghosh, Jenny Chim, Andrew Tran, Yanan Long, Jennifer Mickel, Usman Gohar, Srishti Yadav, Pawan Sasanka Ammanamanchi, Mowafak Allaham, Hossein A. Rahmani, Mubashara Akhtar, Felix Friedrich, Robert Scholz, Michael Alexander Riegler, Jan Batzner, Eliya Habba, Arushi Saxena, Anastassia Kornilova, Kevin Wei, Prajna Soni, Yohan Mathew, Kevin Klyman, Jeba Sania, Subramanyam Sahoo, Olivia Beyer Bruvik, Pouya Sadeghi, Sujata Goswami, Angelina Wang, Yacine Jernite, Zeerak Talat, Stella Biderman, Mykel Kochenderfer, Sanmi Koyejo, and Irene Solaiman. 2025. Who Evaluates AI’s Social Impacts? Mapping Coverage and Gaps in First- and Third-Party Evaluations. *arXiv preprint arXiv:2511.05613* (2025). Version 1, cs.CY.
 - [31] Paul Röttger, Bertie Vidgen, Dong Nguyen, Zeerak Waseem, Helen Margetts, and Janet Pierrehumbert. 2021. HateCheck: Functional Tests for Hate Speech Detection Models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)* (Online). Association for Computational Linguistics. doi:10.18653/v1/2021.acl-long.4
 - [32] Maarten Sap, Dallas Card, Saadia Gabriel, Yejin Choi, and Noah A. Smith. 2019. The Risk of Racial Bias in Hate Speech Detection. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, Anna Korhonen, David Traum, and Lluís Màrquez (Eds.). Association for Computational Linguistics, Florence, Italy, 1668–1678. doi:10.18653/v1/P19-1163
 - [33] Maarten Sap, Saadia Gabriel, Lianhui Qin, Dan Jurafsky, Noah A. Smith, and Yejin Choi. 2020. Social Bias Frames: Reasoning about Social and Power Implications of Language. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel Tetreault (Eds.). Association for Computational Linguistics, Online, 5477–5490. <https://aclanthology.org/2020.acl-main.486>
 - [34] Burr Settles. 2009. *Active Learning Literature Survey*. Computer Sciences Technical Report 1648, University of Wisconsin–Madison.
 - [35] Hong Shen, Alicia DeVos, Motahhare Eslami, and Kenneth Holstein. 2021. Everyday Algorithm Auditing: Understanding the Power of Everyday Users in Surfacing Harmful Algorithmic Behaviors. *Proc. ACM Hum.-Comput. Interact.* 5, CSCW2, Article 433 (Oct. 2021), 29 pages. doi:10.1145/3479577
 - [36] Harvineet Singh, Fan Xia, Mi-Ok Kim, Romain Pirracchio, Rumi Chunara, and Jean Feng. 2023. A Brief Tutorial on Sample Size Calculations for Fairness Audits. arXiv:2312.04745 [stat.AP] <https://arxiv.org/abs/2312.04745>
 - [37] Briana Vecchione, Karen Levy, and Solon Barocas. 2021. Algorithmic Auditing and Social Justice: Lessons from the History of Audit Studies. In *Equity and Access in Algorithms, Mechanisms, and Optimization* (– NY USA, 2021-10-05). ACM, 1–9. doi:10.1145/3465416.3483294
 - [38] Tom Yan and Chicheng Zhang. 2022. Active fairness auditing. In *Proceedings of the 39th International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 162)*, Kamalika Chaudhuri, Stefanie Jegelka, Le Song, Csaba Szepesvari, Gang Niu, and Sivan Sabato (Eds.). PMLR, 24929–24962. <https://proceedings.mlr.press/v162/yan22c.html>
 - [39] Meg Young, Lassana Magassa, and Batya Friedman. 2019. Toward Inclusive Tech Policy Design: A Method for Underrepresented Voices to Strengthen Tech Policy Documents. *Ethics and Information Technology* 21, 2 (2019), 89–103. doi:10.1007/s10676-019-09497-z
 - [40] Juliette Zaccour, Reuben Binns, and Luc Rocher. 2025. Access Denied: Meaningful Data Access for Quantitative Algorithm Audits. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems (CHI '25)*. Association for Computing Machinery, New York, NY, USA, Article 388, 31 pages. doi:10.1145/3706598.3713963