Trust in Practice: Participatory Red Teaming for Safer and Inclusive Generative AI

Lovely-Frances Domingo
lovely@umd.edu
University of Maryland
College Park, MD, USA

María Isabel Magaña Arango
mimagar@umd.edu
University of Maryland
College Park, MD, USA

Sander Schulhoff
sander@learnprompting.org
HackAPrompt and Learn Prompting
San Francisco, CA, USA

Daniel Greene
dgreene1@umd.edu
University of Maryland
College Park, MD, USA

Abstract

We are interested in participating in the **Human Expertise for AI Red-Teaming and Scalable Evaluation (HEARTS) workshop at CHI 2026**. In this **statement of interest**, we provide an overview of our empirical study *Trust in Practice: Exploring Participatory Red-Teaming for Safer, More Inclusive Generative AI* where we explore how AI red teaming competitions, as forms of *public red teaming* (Singh et al., 2025), engage diverse communities in co-designing safe and trustworthy generative AI systems through educational, competition-style events. As part of **our workshop contributions**, we share some of our initial findings, observations, and challenges from our study around **tooling & infrastructure gaps** that support **worker well-being and safety of AI red team practitioners** when *participating in AI red teaming competitions.* **We welcome feedback, questions, and potential areas of collaboration** from other workshop attendees.

CCS Concepts

• **Human-centered computing → Empirical studies in collaborative and social computing**.

Keywords

AI red teaming, participatory AI, trustworthy AI

ACM Reference Format:
Lovely-Frances Domingo, María Isabel Magaña Arango, Sander Schulhoff, and Daniel Greene. 2026. Trust in Practice: Participatory Red Teaming for Safer and Inclusive Generative AI. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (CHI '26).* ACM, New York, NY, USA, 5 pages. https://doi.org/10.1145/nnnnnnn.nnnnnnn

1 Introduction

The rapid deployment of generative AI (GAI) systems has led to calls for multi-stakeholder participation in designing, evaluating, and governing AI systems [23, 30, 34, 41]. Multi-stakeholder participation fosters interdisciplinary collaboration and accountability, increasing trust and safety by addressing harms and risks of AI systems in varied, high-stakes scenarios [2, 10, 26]. "Public red teaming" [1] has emerged as a key participatory method in engaging diverse stakeholders in evaluating generative AI systems through "live" red teaming [12] and events designed as AI red teaming competitions [1, 4, 19].

Public red teaming [1, 36] invites the general public to directly engage with the assessment and evaluation of generative AI-produced content, the discovery of AI model's system vulnerabilities, and mapping out of emerging benefits and consequences of using generative AI tools in everyday life [1]. By engaging the public in competition-style events to identify GAI harms & risks, participants gain familiarity with model capabilities and limitations, while contributing to the development of safe and trustworthy GAI systems [1]. *Public red teaming* can help identify security & safety harms and risks overlooked by internal red teams and application system developers [1, 4], enable transparency in GAI model development and foster public trust [14], and can be seen as a form of public participation in broad AI safety & security policy and governance [31].

To broaden discovery and understanding of risks and harms posed by generative AI systems at large, *public red teaming* centers collective experimentation with generative AI models using a combination of tactics and techniques drawn from adversarial machine learning, manipulation and social engineering techniques from cybersecurity when evaluating problematic generative AI model behaviors and harms [1]. Creative experimentation is essential to the practice of applying red teaming tactics and techniques to generative AI models [4, 42]. Collective engagement of diverse stakeholders through *public red teaming* also suggest that both knowledge practices *and* knowledge about GAI harms and risks are co-constructed and co-produced by individuals from varied lived experiences, backgrounds, and expertise. However, ambiguity in participatory methods and evaluation of the public's contributions [13, 18] risks what Sloane et al [2022] calls "participation washing" and hitting "participation ceilings" [35], where public engagement is superficial and fail to grant participants meaningful influence over the design of AI systems.

As forms of *public red teaming*, **AI red teaming competitions** — such as **HackAPrompt 2.0** [19], the **Hack The Future: Generative AI Red Teaming at DEFCON** [4] and **Gray Swan Arena**,[1] — are commonly designed in the style of a **Capture-The-Flag competition (CTF)**. Capture-The-Flag competitions "gamify the AI red

CHI '26, Barcelona, Spain
2026. ACM ISBN 978-x-xxxx-xxxx-x/YYYY/MM
https://doi.org/10.1145/nnnnnnn.nnnnnnn

[1] Gray Swan Arena is an AI Red Teaming competition hosted by the company Gray Swan, which features Blue Team and Red Team challenges. The competition platform is available via https://app.grayswan.ai/arena

teaming process, making critical security skills accessible to anyone who is willing to try their hand at it" [4, 19]. CTFs have been a staple in the cybersecurity community since 1996 and have become a notable learning mechanism for those interested in learning about cybersecurity [5, 6, 9, 37]. CTFs facilitate interactions between novice and expert cybersecurity practitioners, and between members of various cybersecurity domains and communities of practice. Similarly to other capture-the-flag competitions in cybersecurity, AI red teaming competitions may feature tasks and challenges that emulate real-world scenarios in which participants attack AI models and systems in controlled, sandboxed environments and learn from and share with a community of game or CTF enthusiasts, AI/ML/NLP researchers, security researchers, and red teaming practitioners.

Through a **qualitative case study of HackAPrompt 2.0**, a global ***public red teaming*** event, we situate our work among empirical studies of capture-the-flag competitions in cybersecurity and public red teaming events for generative AI [1, 12, 14, 21, 36, 37, 39, 43]. Our work studies a diverse, global cohort (26+ countries) participating in AI red teaming competition [19] to extend the conversation on generative AI red teaming, safety and trust beyond Western-centric perspectives — which dominate the current academic discourse. Building on Gillespie et al [2024], Singh et al. [2025], and Feffer et al. [2023], we examine sociotechnical processes and human factors in *public red teaming* in a diverse global context.

We contribute to current empirical studies about AI red teaming which examine tactics and techniques [16, 25, 28, 29, 40], challenges to accountability, trust, and safety [1, 12, 13, 22, 24], and calls for inclusion of more perspectives particularly from the Global South [13]. Combining stakeholder analysis, participant observations, semi-structured interviews, our study also builds on Magaña & Shilton [2025] to derive a nuanced understanding of factors which shape public engagement and trust. Through an interdisciplinary blend of social science, critical security studies, and AI research, our study generates novel perspectives about the practice of generative AI red teaming, the connection between public participation in AI red teaming and in evaluating and designing safe and trustworthy AI systems.

2 Case Study

This section provides a brief overview of our ongoing study, in collaboration with HackAPrompt 2.0. Our study protocols were reviewed and approved by our institution's IRB (UMD IRB: 2359459-1).

2.1 HackAPrompt 2.0 AI Red Teaming Competition

HackAPrompt 2.0[2] is an AI red teaming competition founded by our collaborator Sander Schulhoff [28, 29] and his team at Learn Prompting [19]. HackAPrompt 2.0 was held from April–December 2025 and offered $100,000 in total prizes available to participants who compete. This past year's sponsors included companies OpenAI, CATO Networks, Pangea, Telus Digital, and Towards AI [19].

HackAPrompt 2.0 was hosted primarily on a dedicated online platform where all *competition challenges* were available to participants. Each competition track was designed to emulate real-world scenarios where AI systems might be exploited — from AI-based assistants that provide advice on the construction of bio weapons to *agentic AI* being corrupted in Slack channels [19]. Organizers, competitors, and volunteers used Discord as their main communication tool throughout the duration of the event.

2.2 Limitations

Although our study features a *global* AI hacking competition, we acknowledge our own positionality as US-based researchers and that the interviews, artifacts, and other associated materials we will examine are primarily in English. Similarly, interview participants who self-select to join our study may be limited to those who are able to communicate and articulate their practices in English. Although two members of our research team are native speakers of two Filipino languages and some Japanese (Domingo), and of Spanish (Magaña), we recognize that there will be words and concepts that have no direct English translations, or lose a degree of meaning when translated in English. This fact may affect our sample population, which we will make note in our analysis.

Additionally, given that our study is still in progress, and with respect to our collaborators and study population (e.g., hackers, AI red team professionals, researchers in countries with strict anti-hacking laws), we are limited in what we are able to report and share out due confidentiality, security, and privacy concerns.

2.3 Methods

Our study adopts a framework based on Hoff & Bashir [2015] and values in design principles [15, 32]. Hoff & Bashir [2015] suggests that **social and environmental factors** *scaffold* people's perceptions and experiences and influence how they develop "willingness of human operators to rely on automation" when interacting with automated systems:

(1) **Dispositional**: Perceptions, techniques, and strategies based on demographic and cultural background
(2) **Situational**: Design of competition, platform, communication channels, and design and complexity of red teaming tasks
(3) **Learned**: Components of trust, tactics, techniques, and experiences which emerge as through and during the competition, e.g., when competitors dissect GAI models and discuss with others

In our study, we apply this framework to describe how dispositional, situational, and learned factors shape participation and the day-to-day lived experiences of participants during the HackAPrompt 2.0 competition. These factors can be in the form of platform design, workflows and processes, sociotechnical configurations, artifacts, and discourse which emerge from the event. Using values in design principles [8, 15, 32], we situate social and environmental "scaffolding" [20] as "affordances" [7, 8] or "levers" [32] which offer cues to possible experiences, interactions, or actions. Our combined framework acknowledges the role of feelings, emotions, and affect as additional determinants of how participants

[2] https://www.hackaprompt.com

experience events, situations, and/or interactions as they perform and fulfill tasks [8, 20, 32].

2.4 Study Goals

2.4.1 Understanding Organizer Values and Motivations. We examine how human centered values shape and inform the design of AI red teaming competition by adapting values in design approach to analyze organizers' values and motivations in the design of the CTF platform and event, while embedding secure AI norms as participants earn points tackling security challenges. We explore how do organizers collect, analyze, and share and gather input from competitors to enhance learned trust and trustworthiness.

2.4.2 Factors Shaping Situational and Learned Trust. We investigate the dynamics of a global community of practitioners participating in ***public red teaming***, shedding light on how diverse, multilingual stakeholders engage in this process. This work is vital for designing inclusive, participatory, and ethically grounded public red teaming initiatives that foster collaboration across cultures and regions. In what way does participating in HackAPrompt provide a better understanding and knowledge of what makes AI systems safe and trustworthy. Our questions also explore how lived experiences and backgrounds influence trust and motivations for participation. How do participants define "trust" and "trustworthiness," secure/not-secure, safe/not-safe AI systems before/after competition? What steps do they take to identify AI risks, and which trust factors (e.g., task complexity, feedback) shape perceptions their perceptions?

2.4.3 AI Red Teaming Practices. Drawing on how practitioners learn and exercise tactics, tools, and techniques during vulnerability and risk discovery (while red teaming and in Capture-The-Flag competitions) [27, 37, 38, 42] we analyze participants' interactions through the competition platform, social norms and discourse, activities, and interactions with other participants. We focus on understanding individual workflows, processes, and decisions and how these could vary depending on their background, experience, and expertise. We ask participants about their perceptions of harms and risks, information seeking behaviors, decision-making practices.

3 Workshop Contributions and Community Support

In this section, we highlight potential workshop contributions and identify areas in which we could find support from colleagues, researchers, and fellow workshop participants.

3.1 Contributions

3.1.1 Bridging tools and infrastructure gaps. Our study offers some insight into event design, stakeholder engagements, and trust factors between organizers and participants. We analyze the protocols used to report harms to determine if they reflect the goals of both designers and participants.

3.1.2 Worker well-being and support. AI red teaming competitions, as a form of ***public red teaming*** [1] provide spaces and opportunities for the public to acquire skills, exposure, community, and monetary compensation. We investigate individual motivations for participation, successes, limitations, and challenges to examine potential areas of opportunities and interventions to address risks to meaningful contributions by the public. Our research also investigates challenges faced by participants themselves whether during competitions or in their individual practice. Through our study, we expand understanding of public red teaming events to include participants from the Global South highlighting perspectives on labor, safety, and well-being are not restricted to Western-centric (US-based) and English language-centric lenses.

3.2 Support from workshop and research community

We welcome feedback, questions, and potential areas of collaboration from other workshop attendees.

4 Research Team

Our team combines expertise and research interests from participatory AI, cybersecurity and AI security, prompt engineering, and future of work.

Lovely-Frances Domingo is a PhD candidate at the University of Maryland College of Information. Her research examines sociomaterial realities of practicing *generative AI red teaming* and of *public participation in generative AI red teaming*. She draws insight from her lived experiences in cybersecurity consulting and community technology to examine implicit and explicit knowledge practices which shape and are shaped by sociocultural understandings of security, safety, harms, and risks. She contributed in empirical studies applying Value Sensitive Design techniques to examine machine learning-based tools supporting volunteer content moderators [3, 11]. She is also a co-PI for UMD PROGRESS-funded study examining community discourses around harms and risks inform adoption of safety technologies in public schools. She received a MS in Information Management from the Information School, and a BA in English with a minor in History from the University of Washington. She is a member of the Ethics & Values in Design (EVID) lab, and a student affiliate of the Values Centered AI (VCAI) Consortium and NSF Institute for Trustworthy AI in Law & Society (TRAILS).

María Isabel Magaña Arango is a Ph.D. candidate in Information Sciences at the University of Maryland. A Fulbright fellow, Magaña examines the transformative impact of technology on society and media, with a particular focus on Latin America. Her research interests encompass Science and Technology Studies (STS) in Latin America, AI governance and participation [24], open government, and transparency. She is a student affiliate of the NSF Institute for Trustworthy AI in Law & Society (TRAILS). She is part of the Research Group in Journalism (GIP) at Universidad de La Sabana (Colombia). She holds a Master's in Investigative Journalism, Data, and Visualization from Rey Juan Carlos University, Spain (2016), and a Master's in Digital Journalism and Communication from La Sabana University, Colombia (2020).

Sander Schulhoff is the founder and CEO of Learn Prompting and HackAPrompt. He is a machine learning researcher and created the first open-source Prompt Engineering guide, which reached 3 million people and taught them to use tools like ChatGPT. Sander also led a team behind *Prompt Report*, the most comprehensive study of prompting ever done [28]. This 76-page survey, co-authored with researchers from the University of Maryland, OpenAI, Microsoft, Google, Princeton, Stanford, and other leading institutions, analyzed 1,500+ academic papers and covered 200+ prompting techniques. He has published in EMNLP, NeurIPS, ICML, and was a ML Alignment and Theory Scholar (MATS) scholar. He received a Bachelor of Science in Computer Science from the University of Maryland.

Daniel Greene, PhD is an Associate Professor in the College of Information at the University of Maryland (UMD). With a background in sociology and organizational studies, his ethnographic, historical, and social-theoretical research investigates the future of work, and who is included in that future. Greene's 2021 book *The Promise of Access: Technology, Inequality, and the Political Economy of Hope* received the McGannon Book Award for the best book of the year concerning media policy, activism, and social justice. He has also received a 2020 Honorable Mention for Best Paper at CHI, and an Honorable Mention for the 2024 David Edge Prize for the best article in science and technology studies. Dr. Greene is currently working with Dr. Tamara Clegg, also of UMD, on a $3.4 million NSF award to combat systemic racism in STEM education & sport through data literacy. His journal articles appear in such venues as Big Data & Society, New Media & Society, Social Studies of Science, and Research in the Sociology of Work.

Acknowledgments

Our work is funded by the Institute for Trustworthy AI in Law and Society (TRAILS), which is supported by the National Science Foundation under Award 2229885. All opinions, findings, conclusions and recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of the National Science Foundation. We want to thank Sander Schulhoff, Fady Yanni, and their team at HackAPrompt for collaborating with us in this research project; as well as the volunteers and competitors at HackAPrompt 2.0 competition for participating in our study.

References

[1] AI Risk and Vulnerability Alliance (ARVA), Ranjit Singh, Borhane Blili-Hamelin, Carol Anderson, Emnet Tafesse, Briana Vecchione, Beth Duckles, and Jacob Metcalf. 2025. *Red-teaming in the Public Interest.* Technical Report.

[2] Long Bai, Xiangfei Liu, and Jiacan Su. 2023. ChatGPT: The cognitive effects on learning and memory. *Brain-X* 1, 3 (Sept. 2023), e30. doi:10.1002/brx2.30

[3] Yang Trista Cao, Lovely-Frances Domingo, Sarah Gilbert, Michelle L. Mazurek, Katie Shilton, and Hal Daumé Iii. 2024. Toxicity Detection is NOT all you Need: Measuring the Gaps to Supporting Volunteer Content Moderators through a User-Centric Method. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen (Eds.). Association for Computational Linguistics, Miami, Florida, USA, 3567–3587. doi:10.18653/v1/2024.emnlp-main.209

[4] Sven Cattell. 2023. Generative Red Team Recap. https://aivillage.org/defcon%2031/generative-recap/

[5] Kevin Chung and Julian Cohen. 2014. Learning obstacles in the capture the flag model. In *2014 USENIX summit on gaming, games, and gamification in security education (3GSE 14)*. USENIX Association, San Diego, CA. https://www.usenix.org/conference/3gse14/summit-program/presentation/chung

[6] C. Cowan, S. Arnold, S. Beattie, C. Wright, and J. Viega. 2003. Defcon Capture the Flag: defending vulnerable code from intense attack. In *Proceedings DARPA Information Survivability Conference and Exposition.* IEEE Comput. Soc, Washington, DC, USA, 120–129. doi:10.1109/DISCEX.2003.1194878

[7] Jenny L. Davis. 2020. *How artifacts afford: the power and politics of everyday things.* The MIT Press, Cambridge, Massachusetts.

[8] Jenny L Davis. 2023. 'Affordances' for Machine Learning. In *2023 ACM Conference on Fairness, Accountability, and Transparency.* ACM, Chicago IL USA, 324–332. doi:10.1145/3593013.3594000

[9] DEFCON. [n. d.]. DEF CON® Hacking Conference - Capture the Flag Archive. https://defcon.org/html/links/dc-ctf.html

[10] Fernando Delgado, Solon Barocas, and Karen Levy. 2022. An Uncommon Task: Participatory Design in Legal AI. *Proceedings of the ACM on Human-Computer Interaction* 6, CSCW1 (March 2022), 1–23. doi:10.1145/3512898

[11] L.F. Domingo, Sarah Gilbert, Yang (Trista) Cao, Hal Daumé III, Michelle L. Mazurek, and Katie Shilton. 2026. Care, wisdom, and civics: Value sensitive design of large language model support for online moderation. *ACM Trans. Comput.-Hum. Interact.* (Feb. 2026). doi:10.1145/3796227 Just Accepted.

[12] Michael Feffer, Anusha Sinha, Wesley Hanwen Deng, Zachary C. Lipton, and Hoda Heidari. 2024. Red-Teaming for Generative AI: Silver Bullet or Security Theater? http://arxiv.org/abs/2401.15897 arXiv:2401.15897 [cs].

[13] Michael Feffer, Michael Skirpan, Zachary Lipton, and Hoda Heidari. 2023. From Preference Elicitation to Participatory ML: A Critical Survey & Guidelines for Future Research. In *Proceedings of the 2023 AAAI/ACM Conference on AI, Ethics, and Society.* ACM, Montr\'{e}al QC Canada, 38–48. doi:10.1145/3600211.3604661

[14] Sorelle Friedler, Ranjit Singh, Borhane Blili-Hamelin, Jacob Metcalf, and Brian J Chen. 2023. AI Red-Teaming Is Not a One-Stop Solution to AI Harms:. (2023).

[15] Batya Friedman and David Hendry. 2019. *Value sensitive design: shaping technology with moral imagination.* The MIT Press, Cambridge, Massachusetts.

[16] Deep Ganguli, Liane Lovitt, Jackson Kernion, Amanda Askell, Yuntao Bai, Saurav Kadavath, Ben Mann, Ethan Perez, Nicholas Schiefer, Kamal Ndousse, Andy Jones, Sam Bowman, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Nelson Elhage, Sheer El-Showk, Stanislav Fort, Zac Hatfield-Dodds, Tom Henighan, Danny Hernandez, Tristan Hume, Josh Jacobson, Scott Johnston, Shauna Kravec, Catherine Olsson, Sam Ringer, Eli Tran-Johnson, Dario Amodei, Tom Brown, Nicholas Joseph, Sam McCandlish, Chris Olah, Jared Kaplan, and Jack Clark. 2022. Red Teaming Language Models to Reduce Harms: Methods, Scaling Behaviors, and Lessons Learned. http://arxiv.org/abs/2209.07858 arXiv:2209.07858 [cs].

[17] Tarleton Gillespie, Ryland Shaw, Mary L. Gray, and Jina Suh. 2024. AI Red-Teaming is a Sociotechnical System. Now What? doi:10.48550/arXiv.2412.09751 arXiv:2412.09751 [cs] version: 1.

[18] Lara Groves, Aidan Peppin, Andrew Strait, and Jenny Brennan. 2023. Going public: the role of public participation approaches in commercial AI labs. In *2023 ACM Conference on Fairness Accountability and Transparency.* ACM, Chicago IL USA, 1162–1173. doi:10.1145/3593013.3594071

[19] HackAPrompt. 2025. HackAPrompt 2.0. https://hackaprompt.com

[20] Kevin Anthony Hoff and Masooda Bashir. 2015. Trust in Automation: Integrating Empirical Evidence on Factors That Influence Trust. *Human Factors: The Journal of the Human Factors and Ergonomics Society* 57, 3 (May 2015), 407–434. doi:10.1177/0018720814547570

[21] Wm Matthew Kennedy, Cigdem Patlak, Jayraj Dave, Blake Chambers, Aayush Dhanotiya, Darshini Ramiah, Reva Schwartz, Jack Hagen, Akash Kundu, Mouni Pendharkar, Liam Baisley, Theodora Skeadas, and Rumman Chowdhury. 2025. Ask What Your Country Can Do For You: Towards a Public Red Teaming Model. doi:10.48550/arXiv.2510.20061 arXiv:2510.20061 [cs].

[22] Shayne Longpre, Sayash Kapoor, Kevin Klyman, Ashwin Ramaswami, Rishi Bommasani, Borhane Blili-Hamelin, Yangsibo Huang, Aviya Skowron, Zheng-Xin Yong, Suhas Kotha, Yi Zeng, Weiyan Shi, Xianjun Yang, Reid Southen, Alexander Robey, Patrick Chao, Diyi Yang, Ruoxi Jia, Daniel Kang, Sandy Pentland, Arvind Narayanan, Percy Liang, and Peter Henderson. 2024. A Safe Harbor for AI Evaluation and Red Teaming. http://arxiv.org/abs/2403.04893 arXiv:2403.04893 [cs].

[23] Teresa Lynch and Shirley Gregor. 2004. User participation in decision support systems development: Influencing system outcomes. *European Journal of Information Systems* 13, 4 (Dec. 2004), 286–301. doi:10.1057/palgrave.ejis.3000512

[24] Maria Isabel Magaña and Katie Shilton. 2025. Frameworks, Methods and Shared Tasks: Connecting Participatory AI to Trustworthy AI Through a Systematic Review of Global Projects. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency (FAccT '25).* Association for Computing Machinery, New York, NY, USA, 2166–2179. doi:10.1145/3715275.3732148

[25] Ethan Perez, Saffron Huang, Francis Song, Trevor Cai, Roman Ring, John Aslanides, Amelia Glaese, Nat McAleese, and Geoffrey Irving. 2022. Red Teaming Language Models with Language Models. http://arxiv.org/abs/2202.03286 arXiv:2202.03286 [cs].

[26] Soaad Qahha and Syed Ishtiaque Ahmed. 2021. Towards a New Participatory Approach for Designing Artificial Intelligence and Data-Driven Technologies. (2021).

[27] Ryan Ellis and Yuan Stevens. 2022. *Bounty Everything: Hackers and the Making of the Global Bug Marketplace.* Technical Report. https://datasociety.net/library/bounty-everything-hackers-and-the-making-of-the-global-bug-marketplace/

[28] Sander Schulhoff, Michael Ilie, Nishant Balepur, Konstantine Kahadze, Amanda Liu, Chenglei Si, Yinheng Li, Aayush Gupta, HyoJung Han, Sevien Schulhoff, Pranav Sandeep Dulepet, Saurav Vidyadhara, Dayeon Ki, Sweta Agrawal, Chau Pham, Gerson Kroiz, Feileen Li, Hudson Tao, Ashay Srivastava, Hevander Da Costa, Saloni Gupta, Megan L. Rogers, Inna Goncearenco, Giuseppe Sarli, Igor Galynker, Denis Peskoff, Marine Carpuat, Jules White, Shyamal Anadkat, Alexander Hoyle, and Philip Resnik. 2025. The Prompt Report: A Systematic Survey of

Prompt Engineering Techniques. doi:10.48550/arXiv.2406.06608 arXiv:2406.06608 [cs].

[29] Sander Schulhoff, Jeremy Pinto, Anaum Khan, Louis-François Bouchard, Chenglei Si, Svetlina Anati, Valen Tagliabue, Anson Liu Kost, Christopher Carnahan, and Jordan Boyd-Graber. 2024. Ignore This Title and HackAPrompt: Exposing Systemic Vulnerabilities of LLMs through a Global Scale Prompt Hacking Competition. http://arxiv.org/abs/2311.16119 arXiv:2311.16119 [cs].

[30] Andrew D. Selbst, Danah Boyd, Sorelle A. Friedler, Suresh Venkatasubramanian, and Janet Vertesi. 2019. Fairness and Abstraction in Sociotechnical Systems. In *Proceedings of the Conference on Fairness, Accountability, and Transparency*. ACM, Atlanta GA USA, 59–68. doi:10.1145/3287560.3287598

[31] Peter M. Shane. 2012. Cybersecurity policy as if ordinary citizens mattered: The case for public participation in cyber policy making. *Isjlp* 8 (2012), 433.

[32] Katie Shilton. 2013. Values Levers: Building Ethics into Design. *Science, Technology, & Human Values* 38, 3 (May 2013), 374–397. doi:10.1177/0162243912436985

[33] Mona Sloane, Emanuel Moss, Olaitan Awomolo, and Laura Forlano. 2022. Participation Is not a Design Fix for Machine Learning. In *Equity and Access in Algorithms, Mechanisms, and Optimization*. ACM, Arlington VA USA, 1–6. doi:10.1145/3551624.3555285

[34] Philip J. Smith, Roger Beatty, Caroline C. Hayes, Adam Larson, Norman D. Geddes, and Michael C. Dorneich. 2012. Human-Centered Design of Decision-Support Systems. In *The Human–Computer Interaction Handbook* (3 ed.), Julie A. Jacko (Ed.). CRC Press, 589–621. doi:10.1201/b11963-ch-26

[35] Harini Suresh, Emily Tseng, Meg Young, Mary L. Gray, Emma Pierson, and Karen Levy. 2024. Participation in the age of foundation models. In *The 2024 ACM Conference on Fairness Accountability and Transparency*. 1609–1621. doi:10.1145/3630106.3658992 arXiv:2405.19479 [cs].

[36] Victor Storchan, Ravin Kumar, Rumman Chowdhury, Seraphina Goldfarb-Tarrant, and Sven Cattell. 2024. 2024 Generative AI Red Teaming Transparency Report. (2024).

[37] Daniel Votipka, Michelle L. Mazurek, Hongyi Hu, and Bryan Eastes. 2018. Toward a Field Study on the Impact of Hacking Competitions on Secure Development. (2018). https://www.semanticscholar.org/paper/a7022d9dbb30239b65fb20336f9bd5d37275f16b

[38] Daniel Votipka, Rock Stevens, Elissa Redmiles, Jeremy Hu, and Michelle Mazurek. 2018. Hackers vs. Testers: A Comparison of Software Vulnerability Discovery Processes. In *2018 IEEE Symposium on Security and Privacy (SP)*. 374–391. doi:10.1109/SP.2018.00003

[39] Daniel Votipka, Eric Zhang, and Michelle L. Mazurek. 2021. HackEd: A Pedagogical Analysis of Online Vulnerability Discovery Exercises. *2021 IEEE Symposium on Security and Privacy (SP)* null (2021), 1268–1285. doi:10.1109/SP40001.2021.00092

[40] Laura Weidinger, John Mellor, Bernat Guillen Pegueroles, Nahema Marchal, Ravin Kumar, Kristian Lum, Canfer Akbulut, Mark Diaz, Stevie Bergman, Mikel Rodriguez, Verena Rieser, and William Isaac. 2024. STAR: SocioTechnical Approach to Red Teaming Language Models. doi:10.48550/arXiv.2406.11757 arXiv:2406.11757 [cs].

[41] Meg Young, Lassana Magassa, and Batya Friedman. 2019. Toward inclusive tech policy design: a method for underrepresented voices to strengthen tech policy documents. *Ethics and Information Technology* 21, 2 (June 2019), 89–103. doi:10.1007/s10676-019-09497-z

[42] Micah Zenko. 2015. *Red team: how to succeed by thinking like the enemy*. Basic Books, a member of the Perseus Books Group, New York.

[43] Alice Qian Zhang, Ryland Shaw, Jacy Reese Anthis, Ashlee Milton, Emily Tseng, Jina Suh, Lama Ahmad, Ram Shankar Siva Kumar, Julian Posada, Benjamin Shestakofsky, Sarah T. Roberts, and Mary L. Gray. 2024. The Human Factor in AI Red Teaming: Perspectives from Social and Collaborative Computing. In *Companion Publication of the 2024 Conference on Computer-Supported Cooperative Work and Social Computing*. ACM, San Jose Costa Rica, 712–715. doi:10.1145/3678884.3687147